

Fall 2019

# System Software Experiment 2

# File I/O

---

Prof. Jinkyu Jeong ([jinkyu@skku.edu](mailto:jinkyu@skku.edu))

TA – Gyusun Lee ([gyusun.lee@csi.skku.edu](mailto:gyusun.lee@csi.skku.edu))

TA – Jiwon Woo ([jiwon.woo@csi.skku.edu](mailto:jiwon.woo@csi.skku.edu))

Computer Systems and Intelligence Laboratory (<http://csi.skku.edu>)

Sungkyunkwan University

# Contents

- File in Unix
- System calls for File I/O
- Standard I/O functions

# Unix Files

- A Unix **file** is a sequence of  $m$  bytes:
  - $B_0, B_1, \dots, B_k, \dots, B_{m-1}$
- All I/O devices are represented as files:
  - **/dev/sda1** (hard disk partition)
  - **/dev/tty2** (terminal)
    - Ctrl + Alt + F1 ~ F7
- Even the kernel is represented as a file:
  - **/dev/mem** (kernel memory image)
  - **/proc** (kernel data structures)

# Unix File Types

- Regular file
  - Contains arbitrary data
- Directory file
  - A file that contains the names and locations of other files
- Character special and block special files
  - Terminals (character special) and disks (block special)
- FIFO (named pipe)
  - A file type used for inter-process communication
- Socket
  - A file type used for network communication between processes

# Unix I/O

## ■ Characteristics

- The elegant mapping of files to devices allows kernel to export simple interface called Unix I/O
- All input and output is handled in a consistent and uniform way (“byte stream”)

## ■ Basic Unix I/O operations (system calls):

- Opening and closing files
  - **open()** and **close()**
- Changing the **current file position** (seek)
  - **lseek()**
- Reading and writing a file
  - **read()** and **write()**

# Opening Files

- Opening a file informs the kernel that you are getting ready to access that file

```
int fd;    /* file descriptor */
if ((fd = open("/etc/hosts", O_RDONLY)) < 0) {
    perror("open");
    exit(1);
}
```

- Returns a small identifying integer **file descriptor**
  - **fd == -1** indicates that an error occurred
- Each process created by a Unix shell begins life with three open files associated with a terminal:
  - 0: standard input
  - 1: standard output
  - 2: standard error

# Closing Files

- Closing a file informs the kernel that you are finished accessing that file

```
int fd;      /* file descriptor */
int retval; /* return value */

if ((retval = close(fd)) < 0) {
    perror("close");
    exit(1);
}
```

- Moral: Always check return codes, even for seemingly benign functions such as `close()`

# Reading Files

- Reading a file copies bytes from the current file position to memory, and then updates file position.

```
char buf[512];
int fd;      /* file descriptor */
int nbytes;  /* number of bytes read */

/* Open file fd ... */
/* Then read up to 512 bytes from file fd */
if ((nbytes = read(fd, buf, sizeof(buf))) < 0) {
    perror("read");
    exit(1);
}
```

- Returns number of bytes read from file `fd` into `buf`
  - `nbytes < 0` indicates that an error occurred.
  - **short counts** (`nbytes < sizeof(buf)`) are possible and are not errors!



# Writing Files

- Writing a file copies bytes from memory to the current file position, and then updates current file position.

```
char buf[512];
int fd;      /* file descriptor */
int nbytes;  /* number of bytes read */

/* Open the file fd ... */
/* Then write up to 512 bytes from buf to file fd */
if ((nbytes = write(fd, buf, sizeof(buf))) < 0) {
    perror("write");
    exit(1);
}
```

- Returns number of bytes written from buf to file fd .
  - **nbytes < 0** indicates that an error occurred.
  - As with reads, **short counts** are possible and are not errors!

# File Offset

- An offset of an opened file can be set explicitly by calling `lseek()`, `lseek64()`

```
char buf[512];
int fd;          /* file descriptor */
off_t pos;       /* file offset */

/* Get current file offset */
pos = lseek(fd, 0, SEEK_CUR);
/* The file offset is incremented by written bytes */
write(fd, buf, sizeof(buf));
/* Set file position to the first byte of the file */
pos = lseek(fd, 0, SEEK_SET);
```

- Returns the new offset of the file `fd`.
  - **nbytes < 0** indicates that an error occurred.
  - An offset can be set beyond the end of the file.
    - If data is written at that point, a file “hole” is created.

# Unix I/O Example

- Copying standard input to standard output one byte at a time.

```
int main(void)
{
    char c;

    while(read(0, &c, 1) != 0)
        write(1, &c, 1);
    exit(0);
}
```

# Dealing with Short Counts

- Short counts can occur in these situations:
  - Encountering (end-of-file) EOF on reads.
  - Reading text lines from a terminal.
  - Reading and writing network sockets or Unix pipes.
- Short counts does not occur in these situations:
  - Reading from disk files (except for EOF)
  - Writing to disk files.
- How should you deal with short counts in your code?

# Dealing with Short Counts

```
ssize_t rio_readn(int fd, void *usrbuf, size_t n)
{
    size_t nleft = n;
    ssize_t nread;
    char *bufp = usrbuf;

    while (nleft > 0) {
        if ((nread = read(fd, bufp, nleft)) < 0) {
            if (errno == EINTR) /* interrupted by sig handler return */
                nread = 0; /* and call read() again */
            else
                return -1; /* errno set by read() */
        }
        else if (nread == 0)
            break; /* EOF */
        nleft -= nread;
        bufp += nread;
    }
    return (n - nleft); /* return >= 0 */
}
```

# File Metadata

- Data about data, in this case file data.
  - Maintained by kernel, accessed by users with the **stat** and **fstat** functions.

```
/* Metadata returned by the stat and fstat functions */
struct stat {
    dev_t      st_dev;      /* device */
    ino_t      st_ino;      /* inode */
    mode_t     st_mode;     /* protection and file type */
    nlink_t    st_nlink;    /* number of hard links */
    uid_t      st_uid;      /* user ID of owner */
    gid_t      st_gid;      /* group ID of owner */
    dev_t      st_rdev;     /* device type (if inode device) */
    off_t      st_size;     /* total size, in bytes */
    unsigned long st_blksize; /* blocksize for filesystem I/O */
    unsigned long st_blocks; /* number of blocks allocated */
    time_t     st_atime;     /* time of last file access */
    time_t     st_mtime;     /* time of last file modification */
    time_t     st_ctime;     /* time of last inode change */
}; /* statbuf.h included by sys/stat.h */
```

# Accessing File Metadata

```
/* statcheck.c - Querying and manipulating a file's meta data */
```

```
int main (int argc, char **argv)
{
    struct stat st;
    char *type, *readok;

    stat(argv[1], &st);
    if (S_ISREG(st.st_mode)) /* file type */
        type = "regular";
    else if (S_ISDIR(st.st_mode))
        type = "directory";
    else
        type = "other";
    if ((st.st_mode & S_IRUSR)) /* OK to read? */
        readok = "yes";
    else
        readok = "no";

    printf("type: %s, read: %s\n", type, readok);
    exit(0);
}
```

```
bass> ./statcheck statcheck.c
type: regular, read: yes
bass> chmod 000 statcheck.c
bass> ./statcheck statcheck.c
type: regular, read: no
```

# Standard I/O Functions

- The C standard library (libc.so) contains a collection of higher-level **standard I/O** functions
- Examples of standard I/O functions:
  - Opening and closing files (**fopen** and **fclose**)
  - Reading and writing bytes (**fread** and **fwrite**)
  - Reading and writing text lines (**fgets** and **fputs**)
  - Formatted reading and writing (**fscanf** and **fprintf**)



# Standard I/O Streams

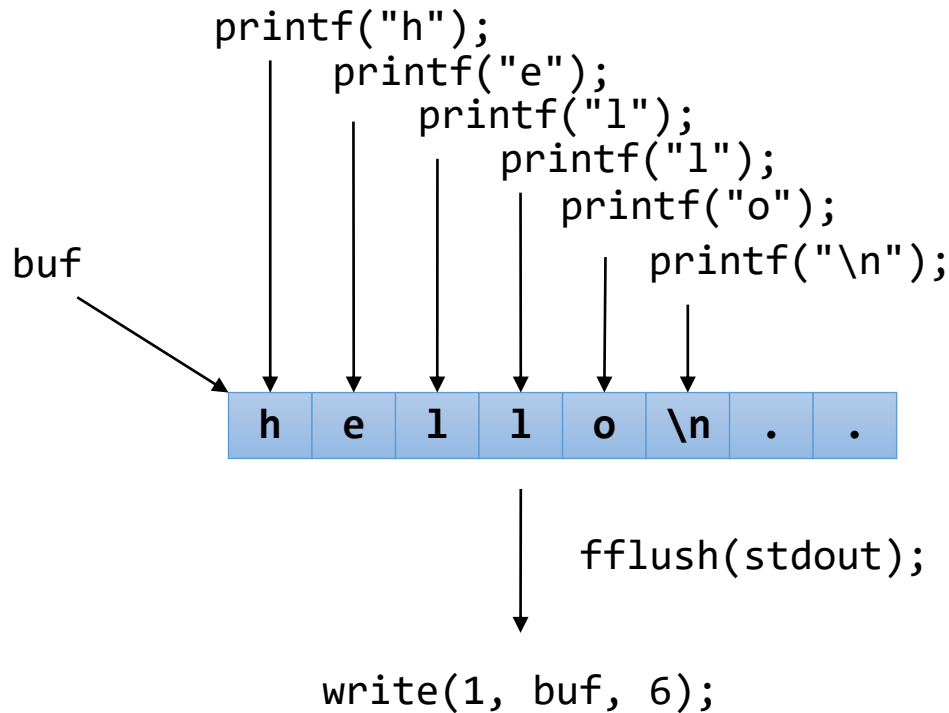
- Standard I/O models open files as **streams**
  - Abstraction for a file descriptor and a buffer in memory
- C programs begin life with three open streams (defined in `stdio.h`)
  - **stdin** (standard input)
  - **stdout** (standard output)
  - **stderr** (standard error)

```
#include <stdio.h>
extern FILE *stdin; /* standard input (descriptor 0) */
extern FILE *stdout; /* standard output (descriptor 1) */
extern FILE *stderr; /* standard error (descriptor 2) */

int main() {
    fprintf(stdout, "Hello, world\n");
}
```

# Buffering in Standard I/O

- Standard I/O functions use buffered I/O



# Buffering in Standard I/O

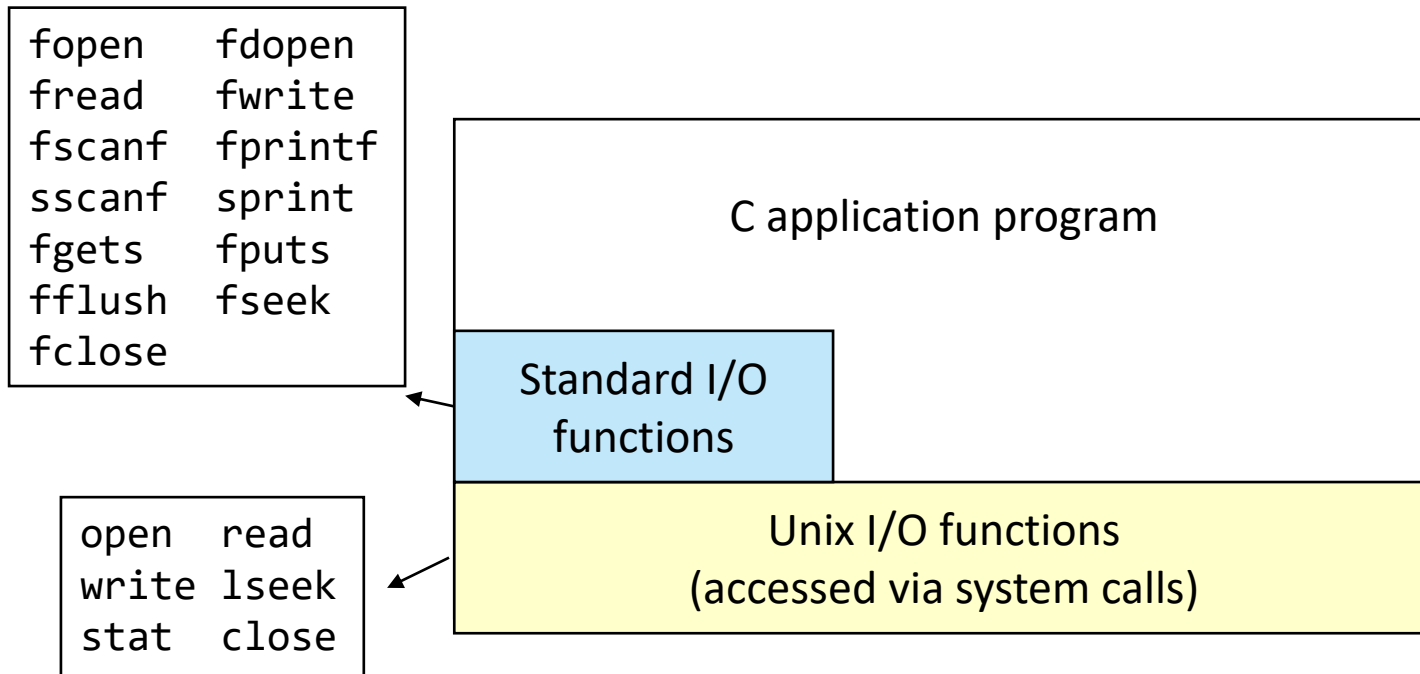
```
#include <stdio.h>

int main()
{
    printf("h");
    printf("e");
    printf("l");
    printf("l");
    printf("o");
    printf("\n");
    fflush(stdout);
    exit(0);
}
```

```
linux> strace ./hello
execve("./hello", ["hello"], [/* ... */]).
...
write(1, "hello\n", 6)                = 6
...
exit_group(0)                         = ?
```

# Unix I/O vs. Standard I/O

- Standard I/O are implemented using low-level Unix I/O



- Which ones should you use in your programs?

# Pros/Cons of Unix I/O

## ■ Pros

- The most general and lowest overhead form of I/O
  - All other I/O packages are implemented on top of Unix I/O functions
- Unix I/O provides functions for accessing file metadata

## ■ Cons

- System call overheads for small-sized I/O
- Dealing with short counts is tricky and error prone
- Efficient reading of text lines requires some form of buffering, also tricky and error prone
- These issues are addressed by the standard I/O

# Pros/Cons of Standard I/O

## ■ Pros

- Buffuring increases efficiency by decreasing the number of **read()** and **write()** system calls
- Shout counts are handled automatically

## ■ Cons

- Provides no function for accessing file metadata
- Standard I/O is not appropriate for input and output on network sockets
  - But there is a way using **fdopen()**

# perror(), errno

- When a system call fails,
  - Returns -1 (or NULL for certain library functions)
  - The latest error information is stored in “errno”.
- perror()
  - Explain error information stored in errno.
  - Print out the information through stderr stream.
- errno
  - Stores int value indicating cause of error.
  - int type extern global variable.
  - Thread-safe!

# errno.h

## ■ /usr/include/asm-generic/errno-base.h

```
1 #ifndef _ASM_GENERIC_ERRNO_BASE_H
2 #define _ASM_GENERIC_ERRNO_BASE_H
3
4 #define EPERM      1 /* Operation not permitted */
5 #define ENOENT     2 /* No such file or directory */
6 #define ESRCH     3 /* No such process */
7 #define EINTR     4 /* Interrupted system call */
8 #define EIO       5 /* I/O error */
9 #define ENXIO     6 /* No such device or address */
10 #define E2BIG     7 /* Argument list too long */
11 #define ENOEXEC   8 /* Exec format error */
12 #define EBADF     9 /* Bad file number */
13 #define ECHILD    10 /* No child processes */
14 #define EAGAIN    11 /* Try again */
15 #define ENOMEM    12 /* Out of memory */
16 #define EACCES    13 /* Permission denied */
17 #define EFAULT    14 /* Bad address */
18 #define ENOTBLK   15 /* Block device required */
19 #define EBUSY     16 /* Device or resource busy */
20 #define EEXIST     17 /* File exists */
21 #define EXDEV     18 /* Cross-device link */
22 #define ENODEV    19 /* No such device */
23 #define ENOTDIR   20 /* Not a directory */
24 #define EISDIR    21 /* Is a directory */
25 #define EINVAL    22 /* Invalid argument */
26 #define ENFILE    23 /* File table overflow */
27 #define EMFILE    24 /* Too many open files */
28 #define ENOTTY    25 /* Not a typewriter */
29 #define ETXTBSY   26 /* Text file busy */
30 #define EFBIG     27 /* File too large */
31 #define ENOSPC    28 /* No space left on device */
32 #define EPIPE     29 /* Illegal seek */
33 #define EROFS     30 /* Read-only file system */
34 #define EMLINK    31 /* Too many links */
35 #define EPIPE     32 /* Broken pipe */
36 #define EDOM      33 /* Math argument out of domain of func */
37 #define ERANGE    34 /* Math result not representable */
38
39 #endif
```

Defined as integers.

Stored in errno when error occurs.

Also accessible through “man errno”



# Handling system call errors

```
#include <stdio.h>

int main()
{
    FILE *fp;
    fp = fopen("file.txt", "r");
    if(fp == NULL){
        perror("Error");
        return (-1);
    }
    fclose(fp);
    return 0;
}
```

```
linux> ./test
linux> Error: No such file or directory
```

# Handling system call errors

```
#include <stdio.h>
```

```
int main()  
{
```

```
    FILE *fp;
```

```
    printf("Hello world!\n");
```

```
    fp = fopen("file.txt", "r");
```

```
    if(fp == NULL){
```

```
        perror("Error");
```

```
        return (-1);
```

```
    }
```

```
    fclose(fp);
```

```
    return 0;
```

```
}
```

stdout

stderr

Send stdout only to result.txt

Hello world!

"re.txt"

```
linux> ./test > result.txt
```

```
linux> Error: No such file or directory
```

# Handling system call errors

```
#include <stdio.h>
```

```
int main()  
{
```

```
    FILE *fp;
```

```
    printf("Hello world!\n");
```

```
    fp = fopen("file.txt", "r");
```

```
    if(fp == NULL){
```

```
        perror("Error");
```

```
        return (-1);
```

```
    }
```

```
    fclose(fp);
```

```
    return 0;
```

```
}
```

stdout

stderr

Send stderr to stdout's dest

```
linux> ./test > result.txt 2>&1
```

```
linux>
```

Error: No such file or directory  
Hello world!

"re.txt"

# Summary

- Unix file I/O
  - `open()`, `read()`, `write()`, `close()`, ...
  - A uniform way to access files, I/O devices, network sockets, kernel data structures, etc.
- When to use standard I/O?
  - When working with disk or terminal files.
- When to use raw Unix I/O
  - When you need to fetch file metadata.
  - When you read or write network sockets or pipes.
  - In rare cases when you need absolute highest performance.

# Remind

- 6 System calls
  - open()
  - close()
  - read()
  - write()
  - lseek()
  - stat() / fstat()

# Example #1 (1)

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```


```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 0  
size: 20



H	e	l	l	o	.	\r	\n
W	h	o		a	r	e	
y	o	u	?				

# Example #1 (2)

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6); // "Hello."
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```


```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 6  
size: 20



H	e	l	l	o	.	\r	\n
W	h	o		a	r	e	
y	o	u	?				

# Example #1 (3)

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2); // "Hello.\r\n"
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```


```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 8  
size: 20



.	e	l	l	o	.	\r	\n
W	h	o		a	r	e	
y	o	u	?				



# Example #1 (4)

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```

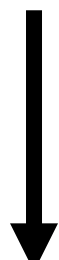
```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 6  
size: 20



H	e	l	l	o	.	\r	\n
W	h	o		a	r	e	
y	o	u	?				

# Example #1 (5)

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```


```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 7  
size: 20



H	e	l	l	o	.	\n	\n
W	h	o		a	r	e	
y	o	u	?				

# Example #1 (6)

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```

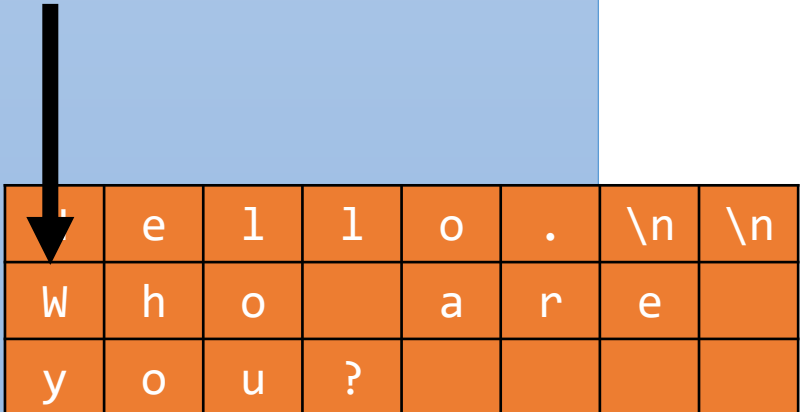
```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 8  
size: 20



h	e	l	l	o	.	\n	\n
W	h	o		a	r	e	
y	o	u	?				

# Example #1 (7)

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```


```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

File state (FD: 3)

path: "hello-dos.txt"  
position: 11  
size: 20



H	e	l		o	.	\n	\n
H	o	w		a	r	e	
y	o	u	?				

# Example #1 (8)

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;
```

File state (FD: 3)  
:CLOSED

```
fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);
```

```
lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);
```

```
lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);
```

```
close(fd);
```

```
fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

# Example #1 (9)

```
char filename[] = "hello-dos.txt";
int fd;
char buffer[16];
off_t pos = 0; // long long;

fd = open(filename, O_RDWR | O_CREAT, 0755);
read(fd, buffer, 6);
read(fd, buffer+6, 2);

lseek(fd, -2, SEEK_CUR);
buffer[0] = '\n';
write(fd, buffer, 1);

lseek(fd, 8, SEEK_SET);
strcpy(buffer, "How");
write(fd, buffer, 3);

close(fd);

fd = open(filename, O_WRONLY | O_CREAT | O_EXCL, 0755);
if (fd < 0)
    printf("errno : %d, error code - EEXIST : %d\n", errno, EEXIST);
```

# Exercise

- Lab exercise #1:
  - Let's make *xtar* utility
  - *tar -cf img.tar 1.jpg 2.jpg*
  - *tar -xf img.tar*
  - It should work same as “tar”
- Download test files on web.
  - `$wget http://csl.skku.edu/uploads/SSE2033F18/lab2.tar.gz`
  - `$tar -xvzf lab2.tar.gz`
  - You can use 2 jpg files for testing your own tar program.

# Exercise

- Handling error cases
  - Not enough input parameters
    - *xtar -cf img.tar*
  - Omitting operation
    - *xtar img.tar 1.jpg 2.jpg*
  - Wrong access
    - *xtar -cf img.tar No\_file.jpg 1.jpg 2.jpg*
  - etc..
- Return error code.
- Print error code & proper error message for each situation.