

# Processes

---

Prof. Jinkyu Jeong ([jinkyu@skku.edu](mailto:jinkyu@skku.edu))

TA – Gyun Lee ([gyusun.lee@csi.skku.edu](mailto:gyusun.lee@csi.skku.edu))

TA – Jiwon Woo ([jiwon.woo@csi.skku.edu](mailto:jiwon.woo@csi.skku.edu))

Computer Systems and Intelligence Laboratory (<http://csi.skku.edu>)

Sung Kyun Kwan University

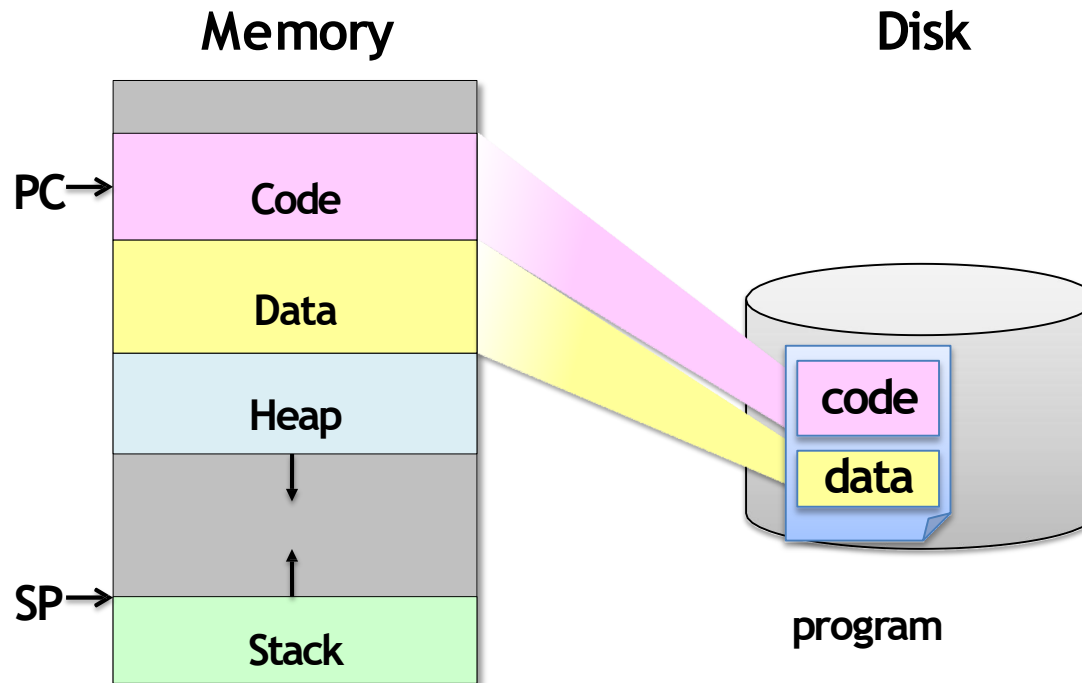
# Processes (1)

- It is not same as “program” or “processor”
- What’s the difference?
  - Program
  - **Process**
  - Processor

# Processes (2)

## ■ Process

- An instance of a program in execution
- One of the most profound ideas in computer science

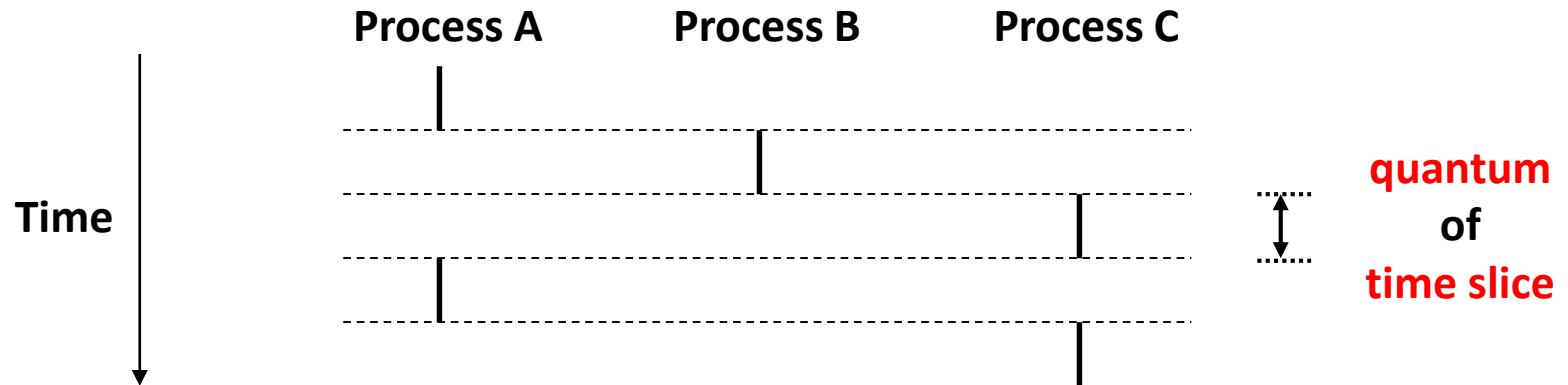


# Processes (3)

- **Process provides each program with two key abstractions:**
  - Logical control flow
    - Each program seems to have exclusive use of the CPU
  - Private address space
    - Each program seems to have exclusive use of main memory
- **How are these illusions maintained?**
  - Process executions interleaved (multitasking).
  - Address space managed by virtual memory system.

# Logical Control Flows

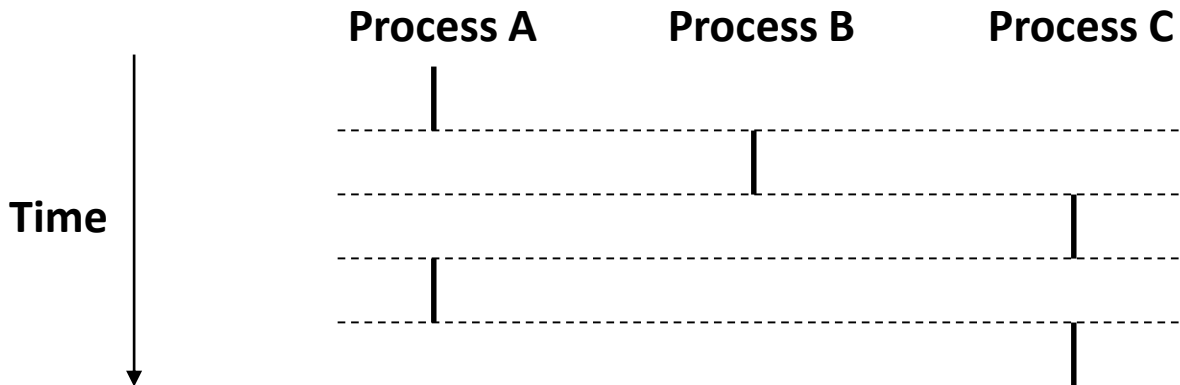
- Each process has its own logical control flow




# Concurrent Processes (1)

## ■ Definition

- Two processes **run concurrently** (are concurrent) if their flows overlap in time.
- Otherwise, they are **sequential**.
- Examples (running on single core):
  - Concurrent: A & B, A & C
  - Sequential: B & C




# Concurrent Processes (2)



# Concurrent Processes (3)

## ■ User View of Concurrent Processes

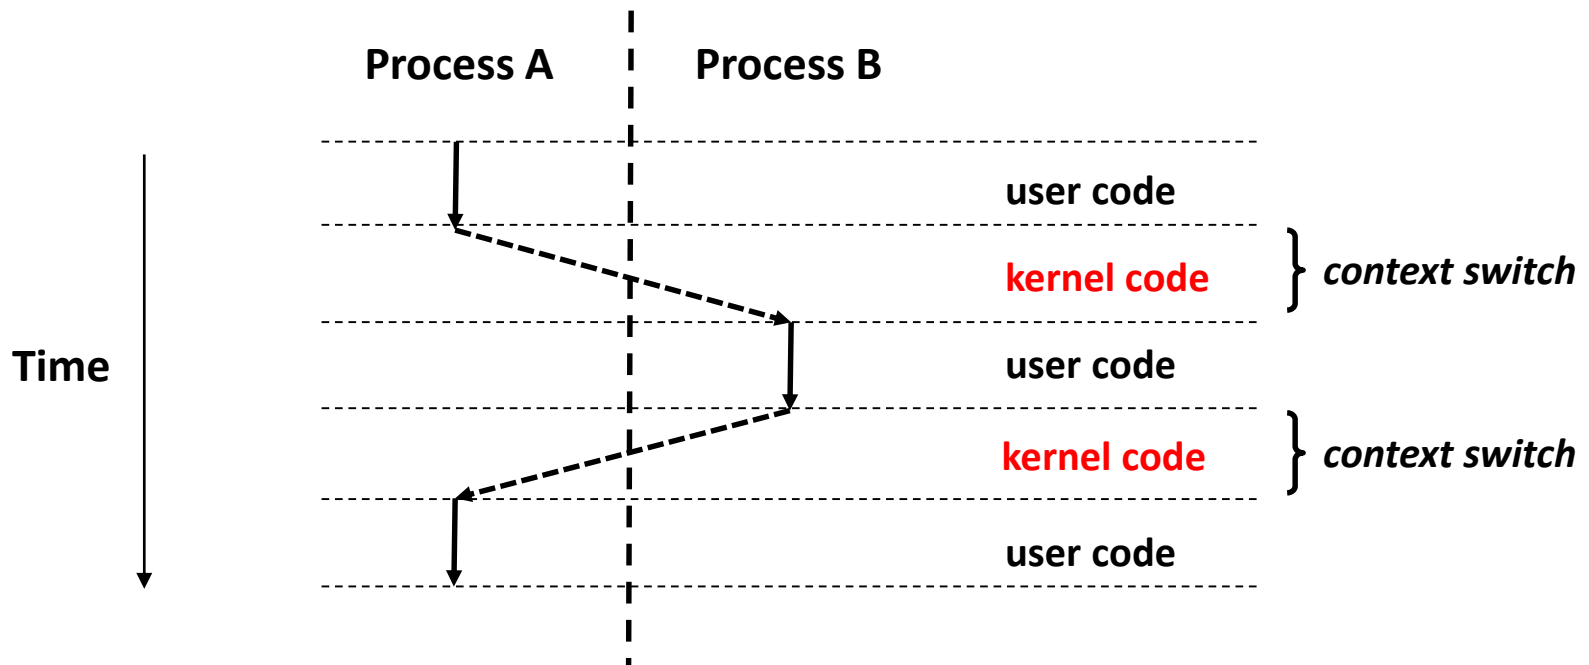
- Control flows for concurrent processes are physically disjoint in time
- However, we can think of concurrent processes are running in parallel with each other





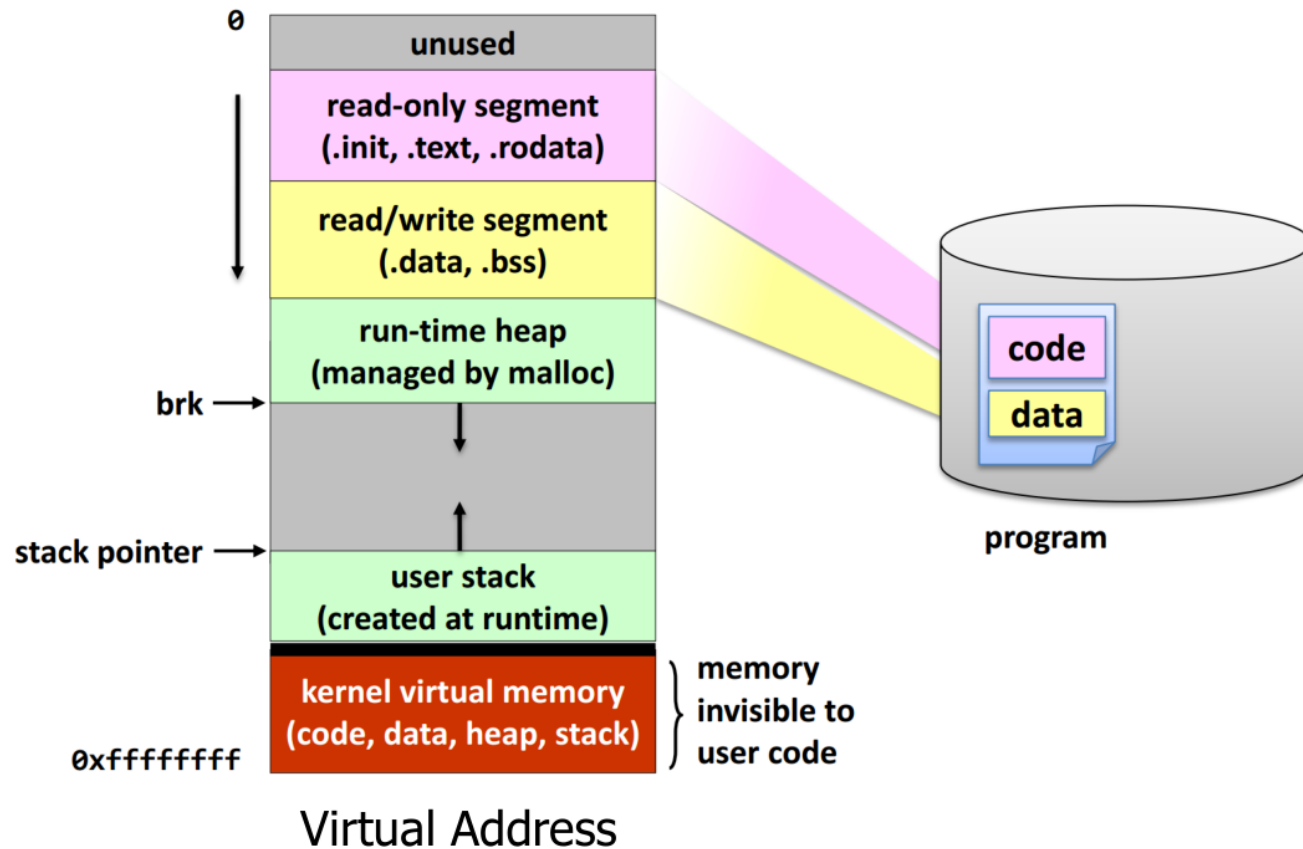
# Context Switching

- Control flow passes from one process to another via a **context switch**



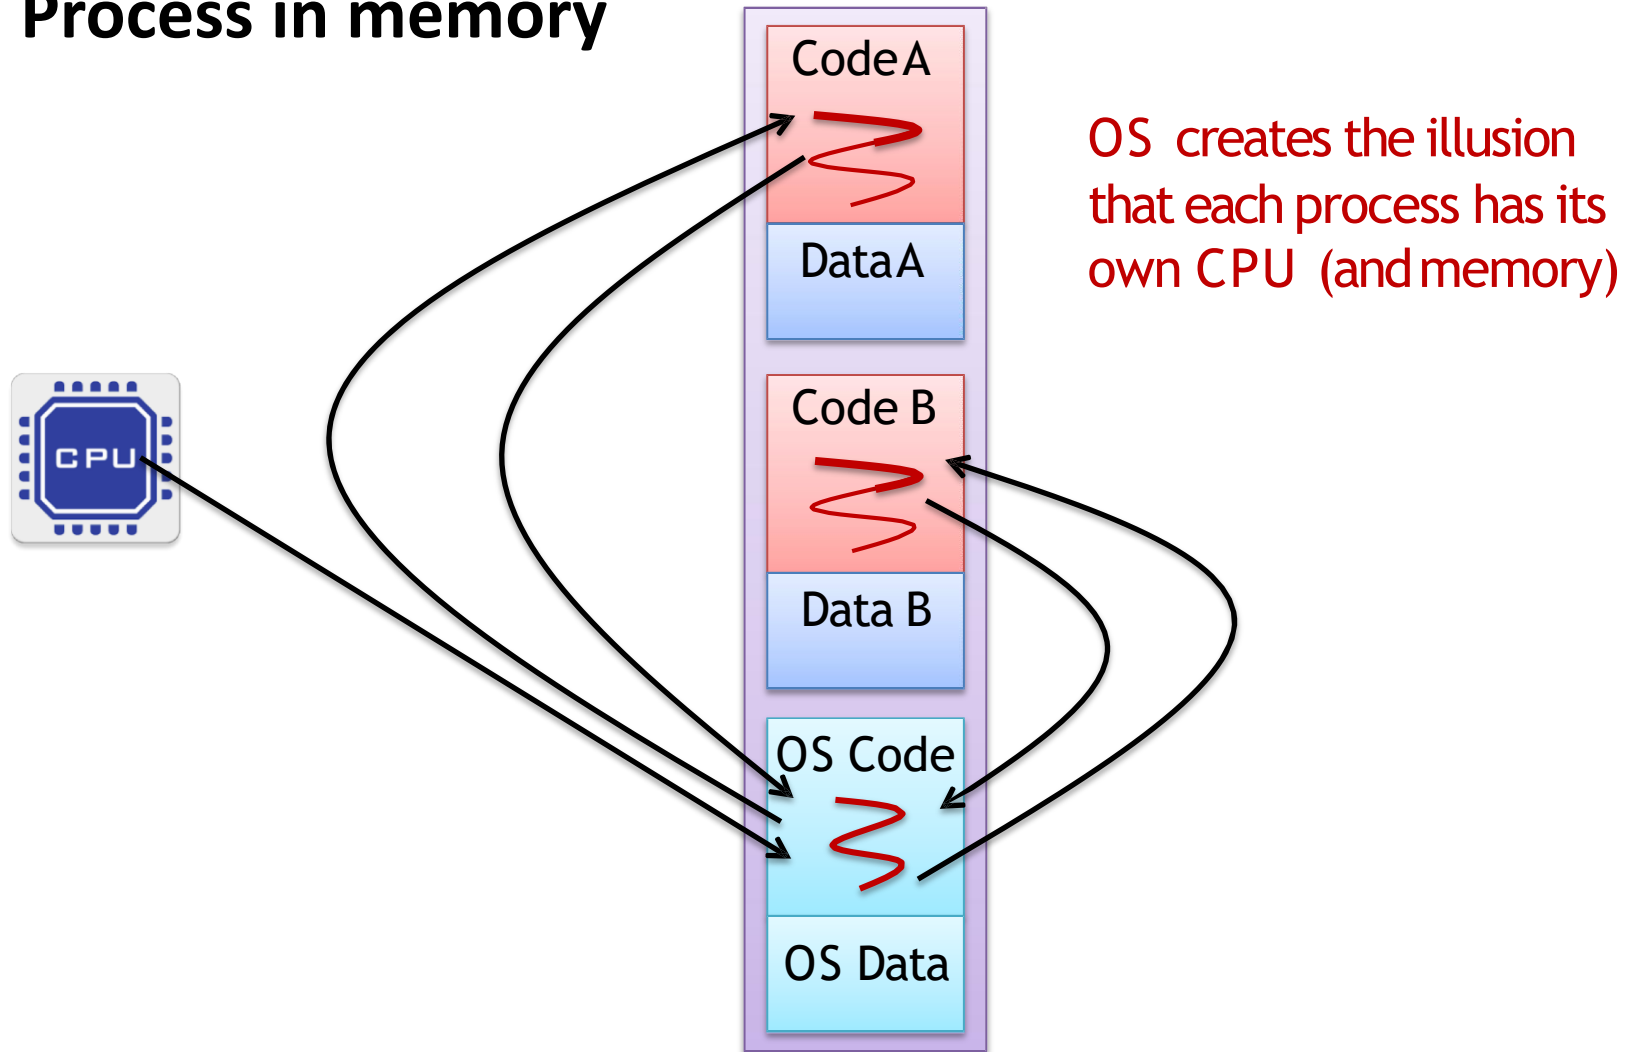
# Private Address Space

## ■ Process in memory

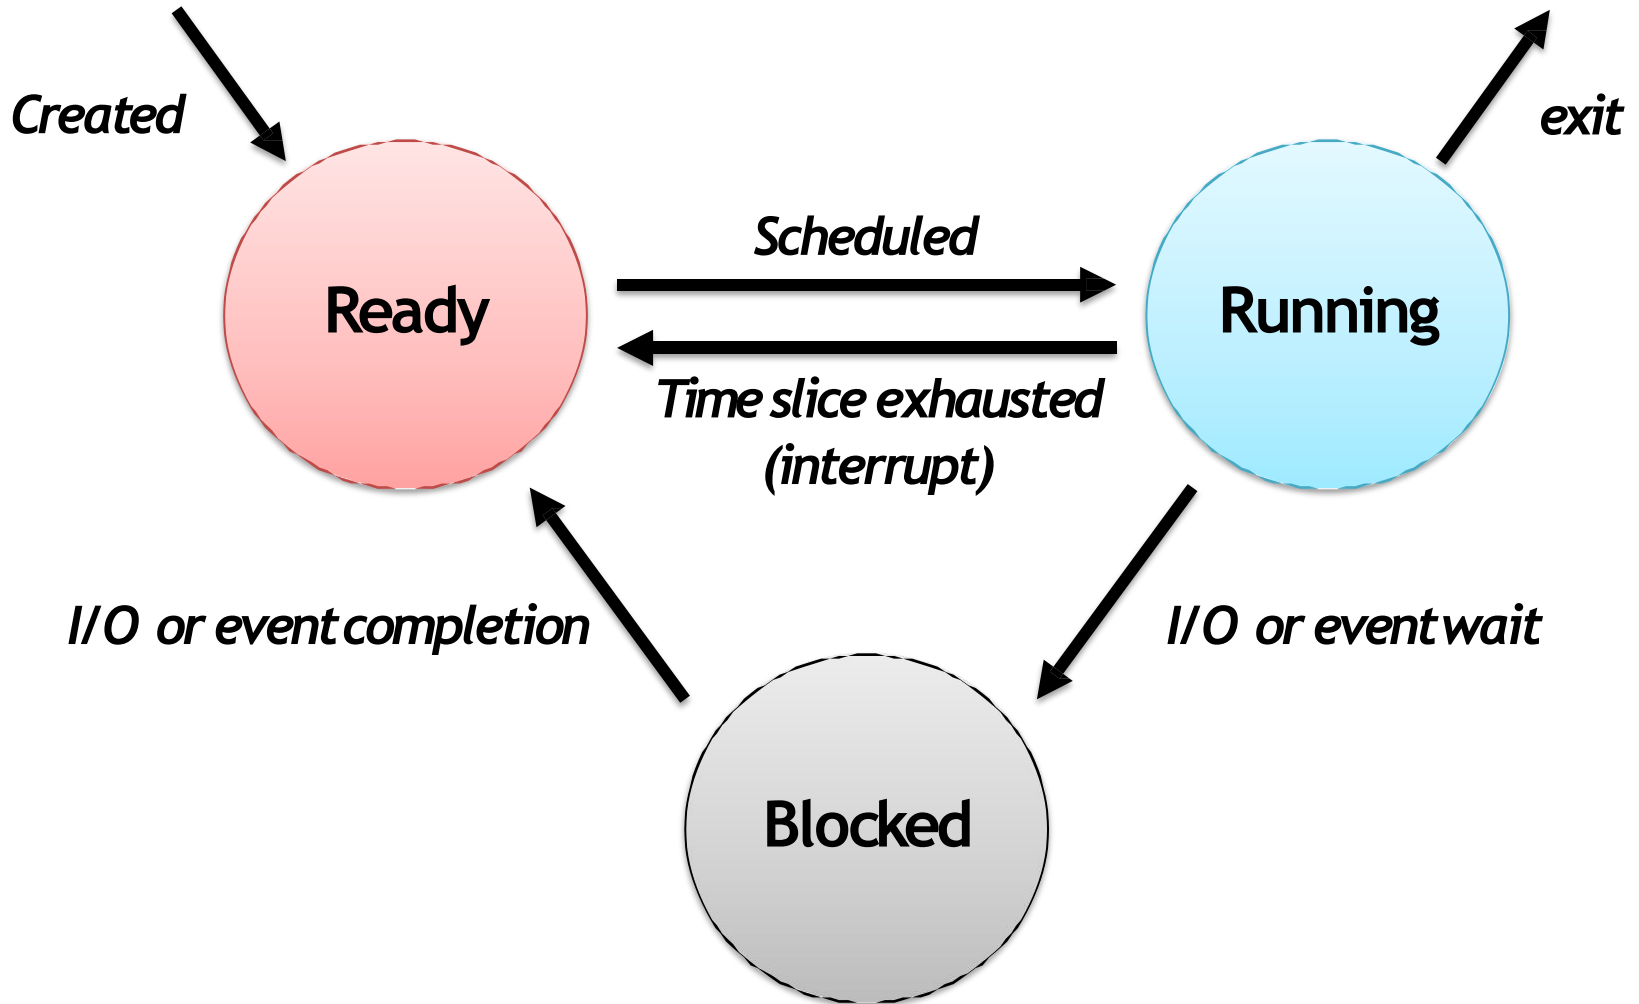


# Private Address Space

## ■ Process in memory



# Process State Transition



# Process Hierarchy


## ■ Parent-child relationship

- One process can create another process
- Unix calls the hierarchy a “process group”
- Windows has no concept of process hierarchy

## ■ Browsing a list of processes:

- ps in Unix
- Task Manager (taskmgr) in Windows

```
$ cat file1 | wc
```



# Creating a New Process

## ■ `pid_t fork(void)`

- Creates a new process (child process) that is identical to the calling process (parent process)
- Returns 0 to the child process
- Returns child's **pid** to the parent process

```
if (fork() == 0) {  
    printf("hello from child\n");  
} else {  
    printf("hello from parent\n");  
}
```

**Fork is interesting  
(and often confusing)  
because it is called  
*once* but returns *twice***

# Fork Example (1)

## ■ Key points

- Parent and child both run **same code**
  - Distinguish parent from child by **return value** from **fork()**
- Start with same state, but each has private copy.
  - Share file descriptors, since child inherits all open files.


```
void fork1() {  
    int x = 1;  
    pid_t pid = fork();  
    if (pid == 0) {  
        printf("Child has x = %d\n", ++x);  
    } else {  
        printf("Parent has x = %d\n", --x);  
    }  
    printf("Bye from process %d with x = %d\n", getpid(), x);  
}
```

# Fork Example (2)

## ■ Key points

- Both parent and child can continue forking.

```
void fork2()
{
    printf("L0\n");
    fork();
    printf("L1\n");
    fork();
    printf("Bye\n");
}
```





# Fork Example (3)

```
#include <sys/types.h>
#include <unistd.h>

int main()
{
    int pid;

    if ((pid = fork()) == 0)
        /* child */
        printf ("Child of %d is %d\n", getppid(), getpid());
    else
        /* parent */
        printf ("I am %d. My child is %d\n", getpid(), pid);
}
```

```
% ./a.out
I am 31098. My child is 31099
Child of 31098 is 31099.
```

```
% ./a.out
Child of 31100 is 31101.
I am 31100. My child is 31101
```

# Process Termination

- **Normal exit (voluntary)**
- **Error exit (voluntary)**
- **Fatal error (involuntary)**
  - Segmentation fault – illegal memory access
  - Protection fault
  - Exceed allocated resources, etc.
- **Killed by another process (involuntary)**
  - By receiving a signal
- **Zombie** process: terminated, but not removed

# Destroying a Process

- **void exit (int status)**

- Exits a process.
  - Normally returns with status 0
- **atexit()** registers functions to be executed upon exit.


```
void cleanup(void) {  
    printf("cleaning up\n");  
}  
  
void fork6() {  
    atexit(cleanup);  
    fork();  
    exit(0);  
}
```

# Synchronizing with Children

- **pid\_t wait (int \*status)**
  - suspends current process until one of its children terminates.
  - return value is the **pid** of the child process that terminated.
  - if **status != NULL**, then the object it points to will be set to a status indicating why the child process terminated.
- **pid\_t waitpid (pid\_t pid, int \*status, int options)**
  - Can wait for specific process
  - Various options

# Wait Example (1)

```
void fork9() {  
    int child_status;  
  
    if (fork() == 0) {  
        printf("HC: hello from child\n");  
    }  
    else {  
        printf("HP: hello from parent\n");  
        wait(&child_status);  
        printf("CT: child has terminated\n");  
    }  
    printf("Bye\n");  
    exit();  
}
```



# Wait Example (2)

- If multiple children completed,
  - will take in arbitrary order.
  - Can use macros **WIFEXITED** and **WEXITSTATUS** to get information about exit status.

```
void fork10() {
    pid_t pid[N];
    int i, child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = 0; i < N; i++) {
        pid_t wpid = wait(&child_status);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```

# Waitpid Example

```
void fork11()
{
    pid_t pid[N];
    int i;
    int child_status;
    for (i = 0; i < N; i++)
        if ((pid[i] = fork()) == 0)
            exit(100+i); /* Child */
    for (i = 0; i < N; i++) {
        pid_t wpid = waitpid(pid[i], &child_status, 0);
        if (WIFEXITED(child_status))
            printf("Child %d terminated with exit status %d\n",
                wpid, WEXITSTATUS(child_status));
        else
            printf("Child %d terminated abnormally\n", wpid);
    }
}
```

# Zombies (1)

## ■ Idea

- When a process terminates, still consumes system resources.
  - Various tables maintained by OS
- Called a “zombie”
  - Living corpse, half alive and half dead

## ■ Reaping

- Performed by parent on terminated child.
- Parent is given exit status information.
- Kernel discards the terminated process.

## ■ What if parent doesn't reap?

- If any parent terminates without reaping a child, then child will be reaped by `init` process.
- Only need explicit reaping for long-running processes.
  - e.g. shells and servers



# Zombies (2)

```
linux> ./forks 7 &
[1] 6639
Running Parent, PID = 6639
Terminating Child, PID = 6640
linux> ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6639 ttyp9        00:00:03 forks
 6640 ttyp9        00:00:00 forks <defunct>
 6641 ttyp9        00:00:00 ps
linux> kill 6639
[1] Terminated
linux> ps
  PID TTY          TIME CMD
 6585 ttyp9        00:00:00 tcsh
 6642 ttyp9        00:00:00 ps
```

```
void fork7()
{
    if (fork() == 0) {
        /* Child */
        printf("Terminating Child, PID = %d\n",
               getpid());
        exit(0);
    } else {
        printf("Running Parent, PID = %d\n",
               getpid());
        while (1); /* Infinite loop */
    }
}
```

- **ps** shows child processes as “defunct”
- Killing parent allows child to be reaped

# Running New Programs (1)

- **int execl (char \*path, char \*arg0, ..., NULL)**
  - loads and runs executable at **path** with arguments **arg0**, **arg1**, ...
    - **path** is the complete path of an executable
    - **arg0** becomes the name of the process
      - Typically **arg0** is either identical to **path**, or else it contains only the executable filename from path.
    - “real” arguments to the executable start with **arg1**, etc.
    - list of args is terminated by a **(char \*) 0** argument.
  - returns **-1** if error, otherwise doesn't return!
- **int execv (char \*path, char \*argv[])**
  - **argv** : null terminated pointer arrays

# Running New Programs (2)

- Example: running /bin/ls

```
main() {  
    if (fork() == 0) {  
        execl("/bin/ls", "ls", "/", NULL);  
    }  
    wait(NULL);  
    printf("completed\n");  
    exit();  
}
```

```
main() {  
    char *args[] = {"ls", "/", NULL};  
    if (fork() == 0) {  
        execv("/bin/ls", args);  
    }  
    wait(NULL);  
}
```

# Summary

## ■ Process abstraction

- Logical control flow
- Private address space

## ■ Process-related system calls

- `fork()`
- `exit()`, `atexit()`
- `wait()`, `waitpid()`
- `execl()`, `execle()`, `execv()`, `execve()`, ...



# Exercise

# Exercise

## ■ Make “xcp” using “cat”

Using dup2 function

- Call “cat <src>” after change fd 1 to <dest>
- \$ xcp <src> <dest>
  - fork()
  - if child
    - close() / open()
    - exec(cat, <src>)
    - return *<last 2 numbers of your student id>*
  - if parent
    - wait(&status)
    - printf(“pid: <pid>”);
    - printf(“status: <status>”);

You should include <fcntl.h>, <unistd.h>,  
<sys/wait.h>, <sys/types.h>  
You can also use <stdio.h>