

Fall 2019

System Software Experiment 2

Pipes and FIFOs

Prof. Jinkyu Jeong (jinkyu@skku.edu)

TA – Gyun Lee (gyusun.lee@csi.skku.edu)

TA – Jiwon Woo (jiwon.woo@csi.skku.edu)

Computer Systems and Intelligence Laboratory (<http://csi.skku.edu>)

Sungkyunkwan University

Contents

- IPC (Inter-Process Communication)
 - Representation of open files in kernel
 - I/O redirection
 - Anonymous Pipe
 - Named Pipe (FIFO)
- Exercise

Everything is a file

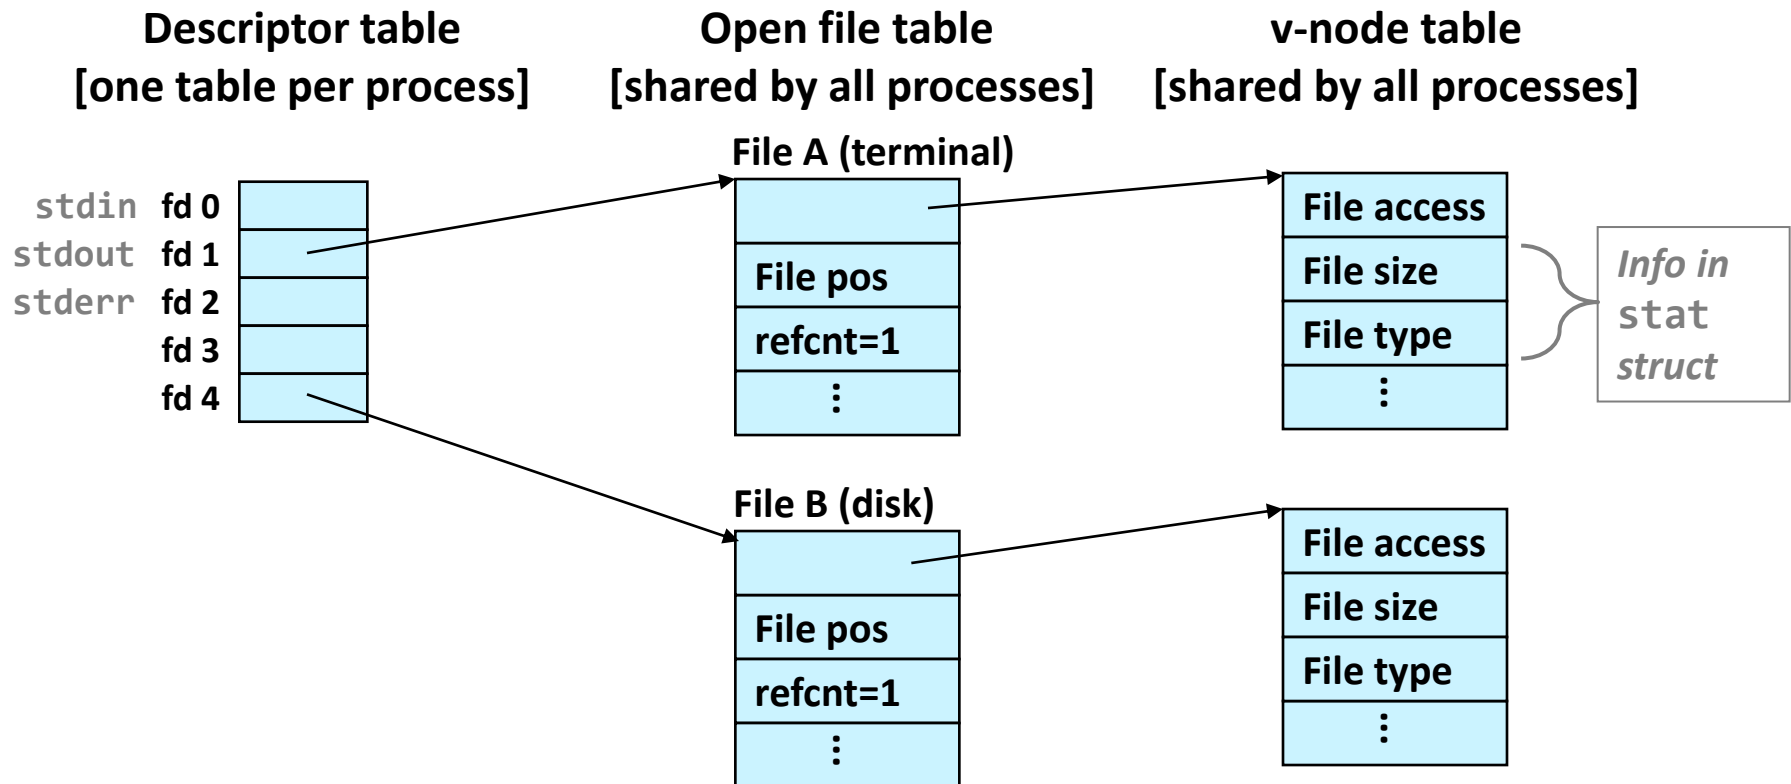
- Actually, “Everything is a file descriptor”
- Pros
 - Can reuse tools, APIs on a wide range of resources
- Cons
 - Not a fast or portable approach
- Communication using file interface?

Open Files in Kernel

- How the Unix kernel represents open files?
- 3-levels
 - Descriptor table
 - 1 table per process
 - Pointer to entry in the “file table”
 - File table
 - Shared by all processes
 - Current file position, mode, reference count, pointer to entry in the “v-node table”
 - v-node table
 - Shared by all processes
 - Information about file itself (size, permission, ...)

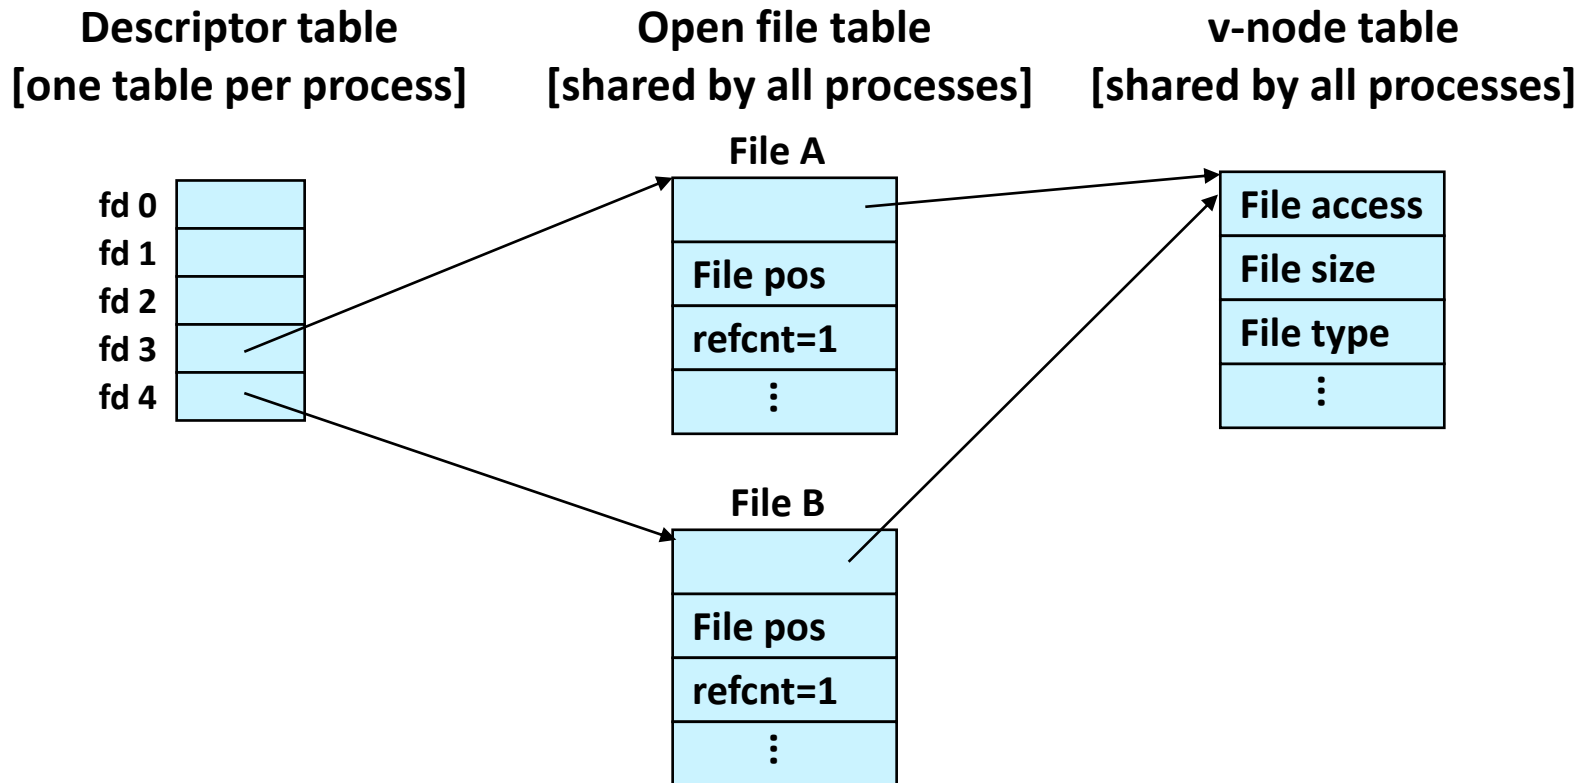
Open Files in Kernel (2)

- How the Unix kernel represents open files?



Open Files in Kernel (3)

- Calling `open()` twice with the same filename



Open Files in Kernel (4)

- Calling fork()

Descriptor table [one table per process] **Open file table** [shared by all processes] **v-node table** [shared by all processes]

Parent's table

fd 0	
fd 1	
fd 2	
fd 3	
fd 4	

Child's table

fd 0	
fd 1	
fd 2	
fd 3	
fd 4	

File A

File pos
refcnt=2
⋮

File B

File pos
refcnt=2
⋮

v-node table

File access
File size
File type
⋮

File access
File size
File type
⋮

Open Files in Kernel (5)

■ What will be the result?

```
int main(void)
{
    char buf[512];
    int fd = open("./tmp.txt", O_RDONLY);

    if (fork() == 0) {
        assert(read(fd, buf, 5) >= 0);
        exit(0);
    } else {
        wait(NULL);
        assert(read(fd, buf, 5) >= 0);
        assert(write(1, buf, 5) >= 0);
        assert(write(1, "\n", 1) >= 0);
    }

    return 0;
}
```

```
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
#include <sys/wait.h>
#include <assert.h>
```


I/O Redirection

- Q: How does a shell implement I/O redirection?

`$ ls > foo.txt`

- A: By calling the `dup2(oldfd, newfd)` function.
 - Copies (per-process) descriptor table entry **oldfd** to entry **newfd**

Descriptor table
before `dup2(4,1)`

fd 0	
fd 1	a
fd 2	
fd 3	
fd 4	b

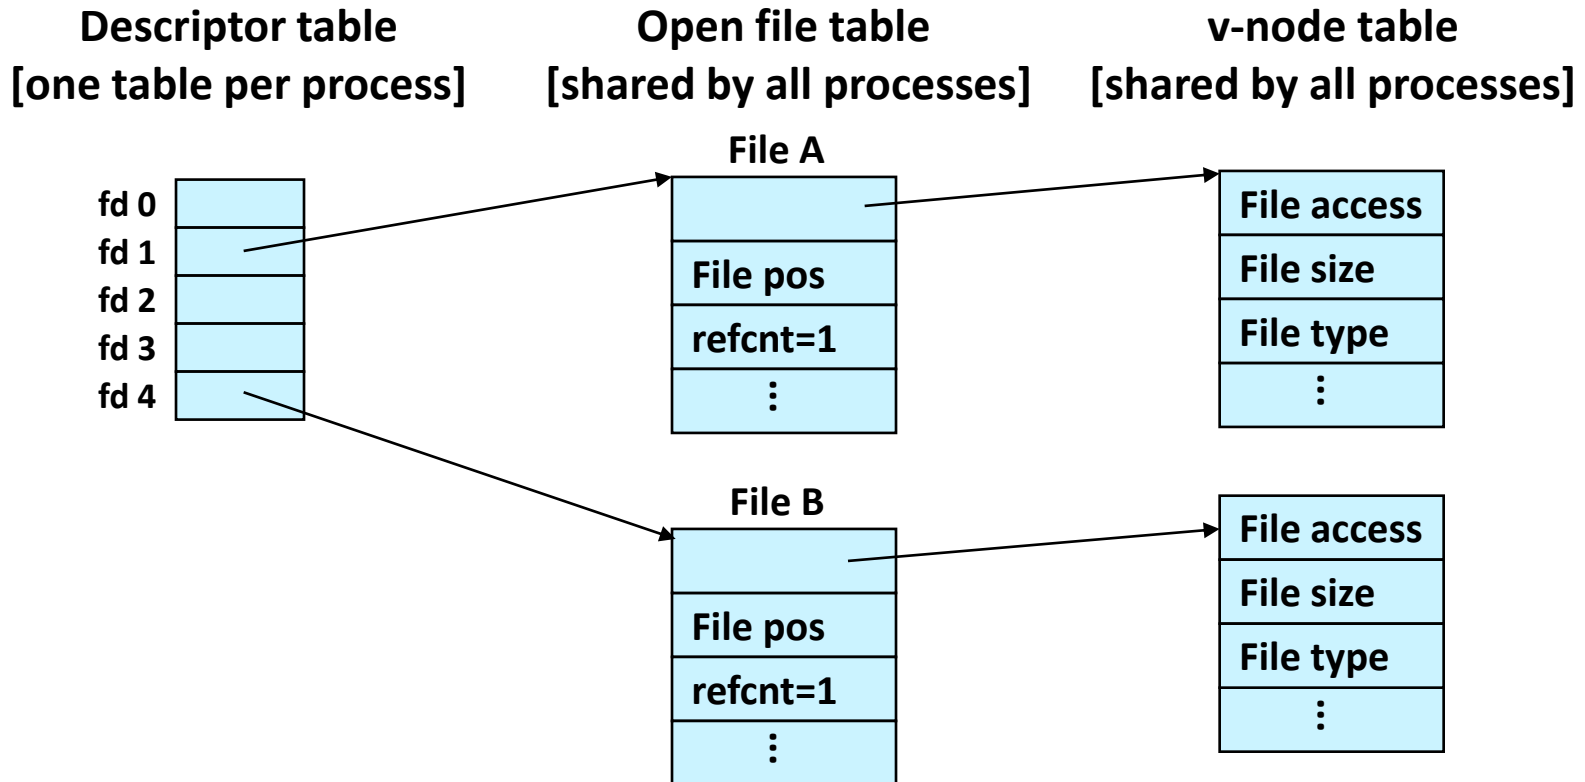


Descriptor table
after `dup2(4,1)`

fd 0	
fd 1	b
fd 2	
fd 3	
fd 4	b

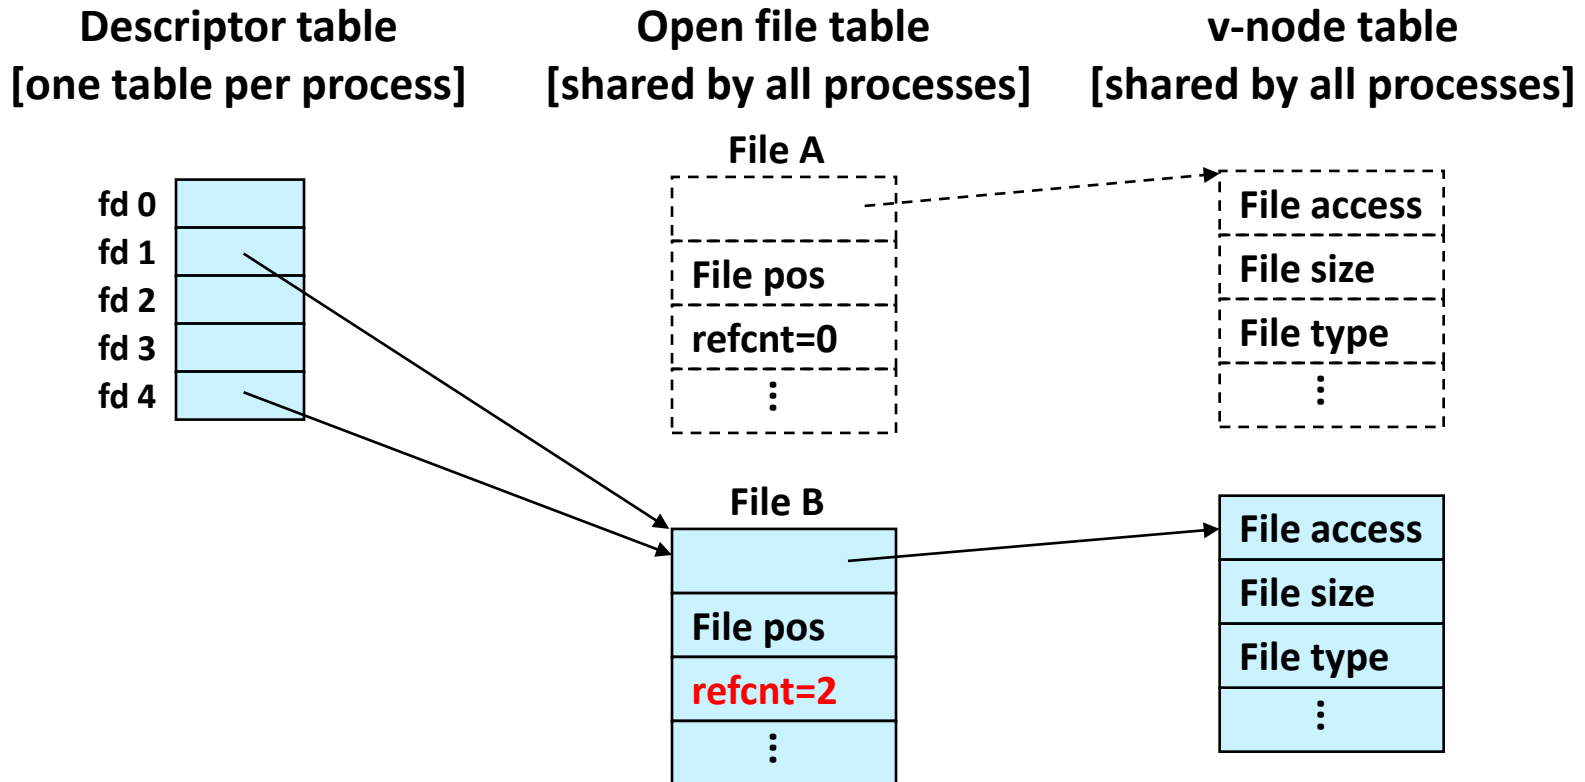
I/O Redirection Example (1)

- Before calling `dup2(4,1)`, `stdout` (descriptor 1) points to a terminal and descriptor 4 points to an open disk file.



I/O Redirection Example (2)

- After calling `dup2(4,1)`, `stdout` is not redirected to the disk file pointed at by descriptor 4.



Pipes

■ Pipes

- The oldest form of UNIX IPC (Inter-process Communication) and provide by all Unix systems.
- IPC using 'file interface'

■ Limitations


- Half-duplex: data flows only in one direction.
- Data only can be read once.

■ Two pipes

- Anonymous pipe
 - No name
- Named pipe
 - We can see it with a file-system

Anonymous Pipe (1)


- `int pipe (int fd[2]);`
 - Two file descriptors are returned through the **fd** argument
 - **fd[0]**: open for reading
 - **fd[1]**: open for writing
 - The output of **fd[1]** is the input for **fd[0]**.



Anonymous Pipe (2)

parent => child:
parent closes fd[0];
child closes fd[1];


parent <= child:
parent closes fd[1];
child closes fd[0];



Anonymous Pipe (3)

parent => child:
parent closes fd[0];
child closes fd[1];

parent <= child:
parent closes fd[1];
child closes fd[0];



Reading/Writing Pipe

- When one end of a pipe is closed,
 - reading from a pipe returns an end of file.
 - writing to a pipe causes **SIGPIPE** is generated and the write returns an error (**EPIPE**).
 - **fstat** function returns a file type of FIFO for the pipe file descriptors (can be tested by **S_ISFIFO** macro)
- You should close unused file descriptors!

Using Anonymous Pipe

```
#include <unistd.h>
#define MAXLINE      80

int main(void)
{
    int n, fd[2];
    pid_t pid;
    char line[MAXLINE];

    if(pipe(fd) < 0) exit(1);
    if((pid = fork()) < 0) exit(2);

    if (pid > 0) {          /* parent */
        close(fd[0]);
        write(fd[1], "hello world\n", 12);
    } else {                /* child */
        close(fd[1]);
        n = read(fd[0], line, MAXLINE);
        write(1, line, n);
    }
    exit(0);
}
```

Named Pipe (FIFO)

- `int mknod (const char *path, mode_t mode, dev_t dev)`
 - `mknod ("path", S_IFIFO, 0);`
- `/usr/bin/mkfifo` program can also be used to make FIFOs on the command line.

Using FIFOs


- Opening a FIFO
 - An open for read(write)-only blocks until some other process opens the FIFO for writing(reading).
- Reading/Writing a FIFO
 - Writing to a FIFO that no process has open for reading causes **SIGPIPE** to generate.
 - When the last writer for a FIFO closes the FIFO, an end of file is generated for the reader of the FIFO.
 - **PIPE_BUF**: the maximum amount of data that can be written atomically to a FIFO (without being interleaved among multiple writers).

Use of FIFOs (1)

■ Duplicating a Stream

- Shell commands to pass data from one shell pipeline to another without creating intermediate temporary files


```
$ mkfifo fifo1  
$ prog3 < fifo1 &  
$ prog1 < infile | tee fifo1 | prog2
```



Use of FIFOs (2)

■ Client-server Communication

- A client-server application to pass data between the client and server **on the same machine**.
 - Clients write to a “well-known” FIFO to send a request to the server.



Summary

- IPC (Inter-Process Communication)
 - Signal
 - Pipe
 - Named pipe (FIFO)
 - Shared memory
 - Semaphore
 - Sockets
 - ...

Exercise

- Make C programs run the following tasks:

```
$ echo "124 * (42 + 3) % 17" | bc
```

- main -> pipe -> fork
 - dup2 -> exec family → echo
 - dup2 -> exec family → bc