

# Pthread

---

Prof. Jinkyu Jeong ([jinkyu@skku.edu](mailto:jinkyu@skku.edu))

TA – Gyun Lee ([gyusun.lee@csi.skku.edu](mailto:gyusun.lee@csi.skku.edu))

TA – Jiwon Woo ([jiwon.woo@csi.skku.edu](mailto:jiwon.woo@csi.skku.edu))

Computer Systems and Intelligence Laboratory (<http://csi.skku.edu>)

Sung Kyun Kwan University

# Creating Concurrent Flows

## ■ Processes

- Kernel automatically interleaves multiple logical flows.
- Each flow has its own private address space.

## ■ Threads

- Kernel automatically interleaves multiple logical flows.
- Each flow shares the same address space.
- Hybrid of processes and I/O multiplexing

## ■ I/O multiplexing with `select()`

- User manually interleaves multiple logical flows
- Each flow shares the same address space
- Popular for high-performance server designs.

**Concurrent Programming**



# **Thread-based**

# Traditional View

- **Process = process context + address space**

## Process context

### Program context:

Data registers

Condition codes

Stack pointer (SP)

Program counter (PC)


### Kernel context:

VM structures

Descriptor table

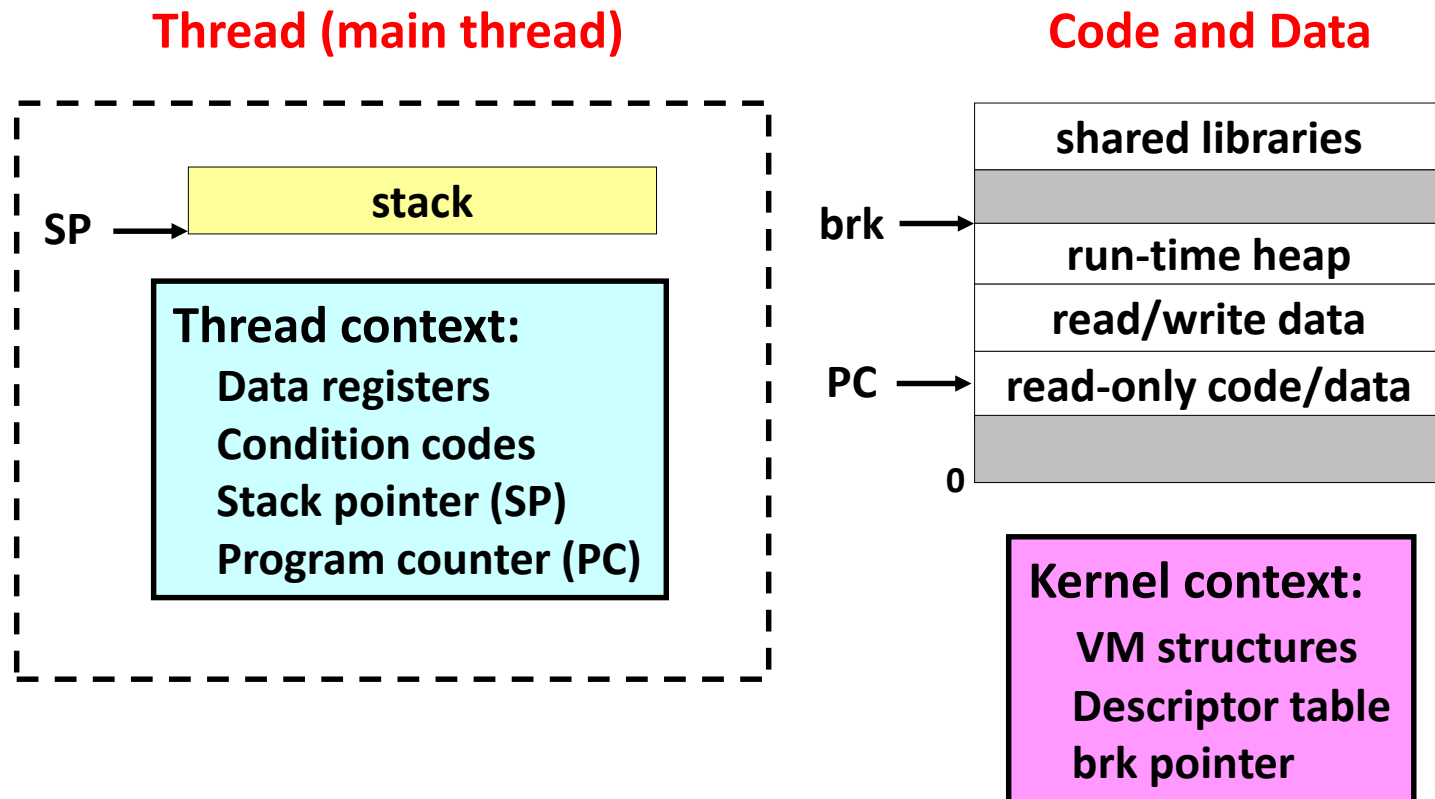
brk pointer

## Code, data, and stack



# Alternate View

- **Process = thread context + kernel context + address space**



# A Process with Multiple Threads

- **Multiple threads can be associated with a process.**
  - Each thread has its own logical control flow (sequence of PC values)
  - Each thread shares the same code, data, and kernel context
  - Each thread has its own thread id (TID)


## Thread 1 (main thread)

stack 1

### Thread 1 context:

Data registers  
Condition codes  
SP1  
PC1

## Shared code and data



## Thread 2 (peer thread)

stack 2


### Thread 2 context:

Data registers  
Condition codes  
SP2  
PC2


# Logical View of Threads

- Threads associated with a process form a pool of peers
  - Unlike processes which form a tree hierarchy

Threads associated with process foo



Process hierarchy



# Threads vs. Processes

## ■ How threads and processes are similar

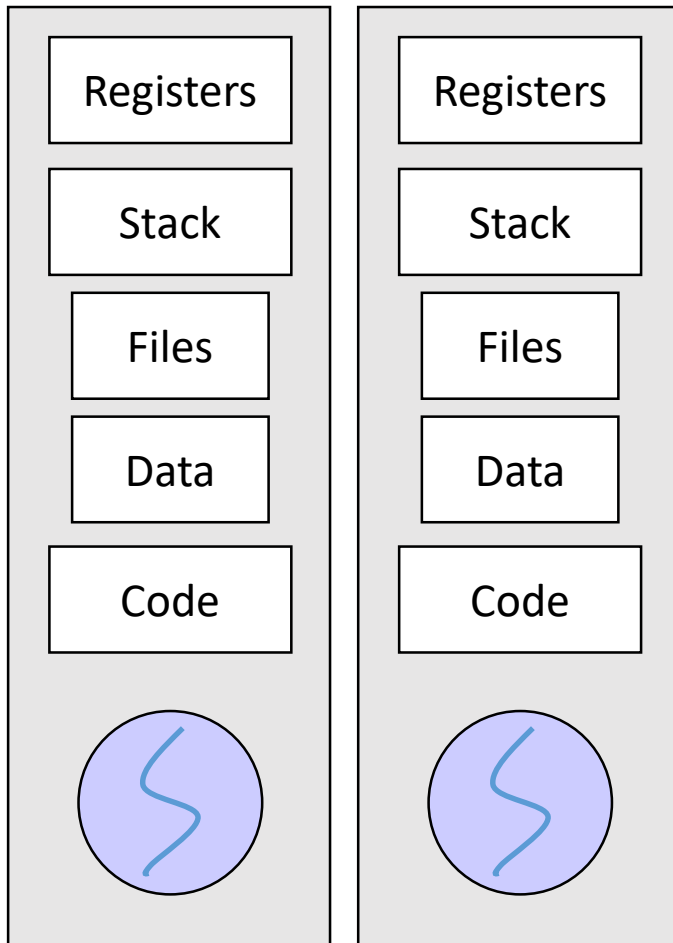
- Each has its own logical control flow.
- Each can run concurrently.
- Each is context switched.

## ■ How threads and processes are different

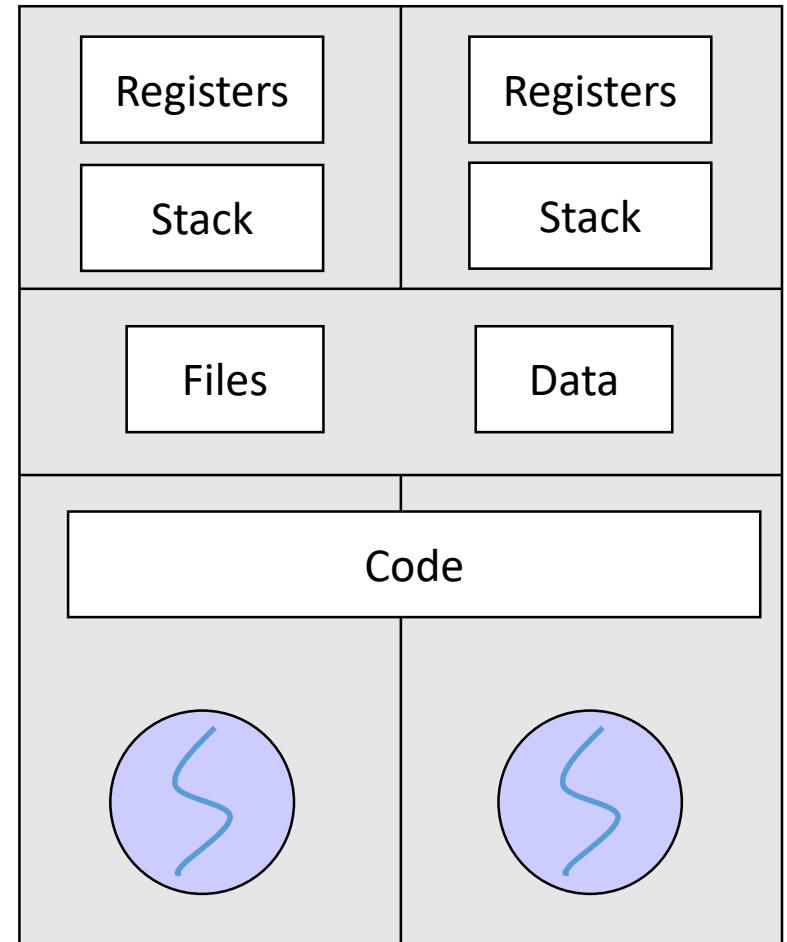

- Threads share code and data, processes (typically) do not.
- Threads are somewhat less expensive than processes.
  - Linux 2.4 Kernel, 512MB RAM, 2 CPUs
    - > 1,811 forks()/second
    - > 227,611 threads/second (125x faster)



# Threads vs. Processes



2 processes



2 threads, 1 process

# The Pthread API

- **ANSI/IEEE POSIX1003.1-1995 Standard**
- **Thread management**
  - Work directly on threads – creating, terminating, joining, etc.
  - Include functions to set/query thread attributes.
- **Mutexes**
  - Provide for creating, destroying, locking and unlocking mutexes.
- **Conditional variables**
  - Include functions to create, destroy, wait and signal based upon specified variable values.

# Pthreads Interface

## ■ POSIX Threads Interface

- Creating and reaping threads
  - `pthread_create()`
  - `pthread_join()`
- Determining your thread ID
  - `pthread_self()`
- Terminating threads
  - `pthread_cancel()`
  - `pthread_exit()`
  - **exit** (terminates all threads), **return** (terminates current thread)
- Synchronizing access to shared variables
  - `pthread_mutex_init()`
  - `pthread_mutex_[un]lock()`
  - `pthread_cond_init()`
  - `pthread_cond_[timed]wait()`
  - `pthread_cond_signal()`, etc.

# Example - process

## Handling Signals (2)

- What is the output of the following program?

```
pid_t pid;
int counter = 2;

void handler1(int sig) {
    counter = counter - 1;
    printf("%d", counter);
    fflush(stdout);
    exit(0);
}

int main() {
    signal(SIGUSR1, handler1);
    printf("%d", counter);
    fflush(stdout);

    if((pid = fork()) == 0) while(1);
    kill(pid, SIGUSR1);
    waitpid(-1, NULL, 0);
    counter = counter + 1;
    printf("%d", counter);
    exit(0);
}
```

21



Computer Systems and Intelligence Lab. - Fall 2019 SSE2033 System Software Experiment 2

5-signals

# Example - thread

```
int gv = 2;
void *thread(void *vargp);
```

```
int main()
{
    pthread_t tid;
    pthread_create(&tid, NULL, thread, NULL);
    pthread_join(tid, NULL);
    printf("[main    ] value:%2d | PID: %5d | TID: %lud\n",
           gv, getpid(), (unsigned)pthread_self() );
    exit(0);
}
```

```
void *thread(void *vargp) /* Thread routine */
{
    printf("Hello, world!\n");
    printf("[created] value:%2d | PID: %5d | TID: %lud\n",
           gv, getpid(), (unsigned)pthread_self() );
    return NULL;
}
```

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/syscall.h>
```

# Creating Threads (1)

- **int pthread\_create** (pthread\_t \*thread, pthread\_attr\_t \*attr, void \*(\*start\_routine)(void \*), void \*arg)
  - **pthread\_create()** returns the new thread ID via the **thread** argument.
    - The caller can use this thread ID to perform various operations on the thread.
  - The **attr** parameter is used to set thread attributes.
    - NULL for the default values.
  - The **start\_routine** denotes the C routine that the thread will execute once it is created.
    - C routine that the thread will execute once it is created.
  - A single argument may be passed to **start\_routine()** via **arg**.

# Creating Threads (2)

## ■ Notes:

- Initially, **main()** comprises a single, default thread.
- All other threads should must be explicitly created by the programmer.
- Once created, threads are peers, and may create other threads.
- The maximum number of threads that may be created by a process is implementation dependent.

# Terminating Threads

- **void pthread\_exit (void \*retval)**
  - **pthread\_exit()** terminates the execution of the calling thread.
    - Typically, this is called after a thread has completed its work and is no longer required to exist.
  - The **retval** argument is the return value of the thread.
    - It can be consulted from another thread using **pthread\_join()**.
  - It does not close files; any files opened inside the thread will remain open after the thread is terminated.



# Cancelling Threads

- **int pthread\_cancel (pthread\_t thread)**
  - **pthread\_cancel()** sends a cancellation request to the thread denoted by the **thread** argument.
  - Depending on its settings, the target thread can then either ignore request, honor it immediately, or defer it till it reaches a cancellation point.
    - pthread\_setcancelstate():  
PTHREAD\_CANCEL\_(ENABLE|DISABLE)
    - pthread\_setcanceltype():  
PTHREAD\_CANCEL\_(DEFERRED|ASYNCHRONOUS)
  - Threads are always created by **pthread\_create()** with cancellation enabled and deferred.

# Joining Threads

- **int pthread\_join (pthread\_t thread, void \*\*retval)**
  - **pthread\_join()** suspends the execution of the calling thread until the thread identified by **thread** terminates, either by calling **pthread\_exit()** or by being cancelled.
  - The return value of **thread** is stored in the location pointed by **retval**.
  - It returns **PTHREAD\_CANCELLED** if thread was cancelled.
  - It is impossible to join a detached thread.

# Detaching Threads

- **int pthread\_detach (pthread\_t thread)**
  - **pthread\_detach()** puts the thread in the detached state.
    - This guarantees that the memory resources consumed by **thread** will be freed immediately when thread terminates.
    - However, this prevents other threads from synchronizing on the termination of thread using **pthread\_join()**.
  - A thread can be detached when it is created:

```
pthread_t tid;  
pthread_attr_t attr;  
  
pthread_attr_init (&attr);  
pthread_attr_setdetachstate(&attr, PTHREAD_CREATE_DETACHED);  
pthread_create(&tid, &attr, start_routine, NULL);  
pthread_attr_destroy (&attr);
```

# Thread Identifiers

- **pthread\_t pthread\_self (void)**
  - **pthread\_self()** returns the unique, system assigned thread ID of the calling thread.
- **int pthread\_equal (pthread\_t t1, pthread\_t t2)**
  - **pthread\_equal()** returns a non-zero value if **t1** and **t2** refer to the same thread.
  - Because thread IDs are opaque objects, the C language equivalence operator **==** should not be used to compare two thread IDs against each other.

# **Threads Synchronization**

# Example

```
#include <stdio.h>
#include <pthread.h>

int num;

void* inc (void* tid) {
    int iter = 10000;
    while(iter--) num++;
}

void* dec (void* tid) {
    int iter = 10000;
    while(iter--) num--;
}

int main()
{
    pthread_t thread_inc, thread_dec;
    pthread_create(&thread_inc, NULL, &inc, NULL);
    pthread_create(&thread_dec, NULL, &dec, NULL);

    pthread_join(thread_inc, NULL);
    pthread_join(thread_dec, NULL);

    printf("%d\n", num);
    return 0;
}
```

# Mutex (1)

- **Mutex is an abbrev. for “mutual exclusion”**
  - Primary means of implementing thread synchronization.
    - Protects shared data when multiple writes occurs.
  - A mutex variable acts like a “lock” protecting access to a shared resource.
    - Only one thread can lock (or own) a mutex variable at any given time.
    - Even if several threads try to lock a mutex, only one thread will be successful. Other threads are blocked until the owner releases the lock.
  - Mutex is used to prevent “race” conditions.
    - race condition: anomalous behavior due to unexpected critical dependence on the relative timing of events.

# Mutex (2)

```
int deposit(int amount)
{
    int balance;

    balance = get_balance();
    balance += amount;
    put_balance(balance);
    return balance;
}
```

```
int withdraw(int amount)
{
    int balance;

    balance = get_balance();
    balance -= amount;
    put_balance(balance);
    return balance;
}
```

**T1 executes deposit(100)**

balance = get\_balance();  
balance += 100;

put\_balance(balance);

**T2 executes withdraw(300)**

balance = get\_balance();  
balance -= 300;  
put\_balance(balance);



# Creating/Destroying Mutexes

## ■ Static initialization

- `pthread_mutex_t m =  
PTHREAD_MUTEX_INITIALIZER;`

## ■ Dynamic initialization

- `pthread_mutex_t m;  
pthread_mutex_init (&m,  
(pthread_mutexattr_t *)NULL);`

## ■ Destroying a mutex

- `pthread_mutex_destroy (&m);`
- Destroys a mutex object, freeing the resources it might hold.

# Using Mutexes (1)

- **int pthread\_mutex\_lock** (pthread\_mutex\_t \*mutex)
  - Acquire a lock on the specified **mutex** variable.
  - If the **mutex** is already locked by another thread, block the calling thread until the **mutex** is unlocked.
- **int pthread\_mutex\_unlock** (pthread\_mutex\_t \*mutex)
  - Unlock a **mutex** if called by the owning thread.
- **int pthread\_mutex\_trylock** (pthread\_mutex\_t \*mutex)
  - Attempt to lock a **mutex**.
  - If the **mutex** is already locked, return immediately with a “busy” error code.

# Using Mutexes (2)

```
pthread_mutex_t m =  
    PTHREAD_MUTEX_INITIALIZER;
```

```
int deposit(int amount)  
{  
    int balance;
```

```
    pthread_mutex_lock(&m);
```

```
    balance = get_balance();  
    balance += amount;  
    put_balance(balance);
```

```
    pthread_mutex_unlock(&m);
```

```
    return balance;  
}
```

```
int withdraw(int amount)  
{  
    int balance;
```

```
    pthread_mutex_lock(&m);
```

```
    balance = get_balance();  
    balance -= amount;  
    put_balance(balance);
```

```
    pthread_mutex_unlock(&m);
```

```
    return balance;  
}
```

# Condition Variables (1)

## ■ Another way for thread synchronization

- While mutexes implement synchronization by controlling thread access to data, condition variables allow threads to synchronize based upon the actual value of data.
- Without condition variables, the programmer would need to have threads continually polling to check if the condition is met.
  - This can be very resource consuming since the thread would be continuously busy in this activity.
- A condition variable is always used in conjunction with a mutex lock.

# Condition Variables (2)

## ■ How condition variables work

- A thread locks a mutex associated with a condition variable.
- The thread tests the condition to see if it can proceed.
- If it can
  - Your thread does its work
  - Your thread unlocks the mutex
- If it cannot
  - The thread sleeps. **The mutex is automatically released.**
  - Some other threads signals the condition variable.
  - Your thread wakes up from waiting **with the mutex automatically locked**, and it does its work.
  - Your thread releases the mutex when it's done.

# Creating/Destroying CV

## ■ Static initialization

- `pthread_cond_t cond =  
    PTHREAD_COND_INITIALIZER;`

## ■ Dynamic initialization

- `pthread_cond_t cond;  
    pthread_cond_init (&cond,  
    (pthread_condattr_t *)NULL);`

## ■ Destroying a condition variable

- `pthread_cond_destroy (&cond);`
- Destroys a condition variable, freeing the resources it might hold.

# Using Condition Variables

- **int pthread\_cond\_wait** (pthread\_cond\_t \*cond, pthread\_mutex\_t \*mutex)
  - Blocks the calling thread until the specified condition is signalled.
  - This should be called while mutex is locked, and it will automatically release the mutex while it waits.
- **int pthread\_cond\_signal** (pthread\_cond\_t \*cond)
  - Signals another thread which is waiting on the condition variable.
  - Calling thread should have a lock.
- **int pthread\_cond\_broadcast**(pthread\_cond\_t \*cond)
  - Used if more than one thread is in a blocking wait state.




# Exercise



# Producer-Consumer

- **Bounded buffer: size N**
- **Producer threads writes data to buffer**
  - Should not write more than N items
- **Consumer threads reads data from buffer**
  - Should not try to consume if there is no data
- **Buffer is circular**



# Producer-Consumer (1)

```
#include <stdio.h>
#include <stdlib.h>
#include <pthread.h>

#define QSIZE          5
#define LOOP           30

typedef struct {
    int data[QSIZE];
    int index;
    int count;
    pthread_mutex_t lock;
    pthread_cond_t notfull;
    pthread_cond_t notempty;
} queue_t;

void *produce (void *args);
void *consume (void *args);
void put_data (queue_t *q, int d);
int  get_data (queue_t *q);
```

# Producer-Consumer (2)

```
int main ()
{
    queue_t *q;
    pthread_t producer, consumer;

    q = qinit();

    pthread_create(&producer, NULL, produce, (void *)q);
    pthread_create(&consumer, NULL, consume, (void *)q);

    pthread_join (producer, NULL);
    pthread_join (consumer, NULL);

    qdelete(q);
}
```

# Producer-Consumer (3)

```
queue_t *qinit()
{
    queue_t *q;

    q = (queue_t *) malloc(sizeof(queue_t));
    q->index = q->count = 0;
    pthread_mutex_init(&q->lock, NULL);
    pthread_cond_init(&q->notfull, NULL);
    pthread_cond_init(&q->notempty, NULL);

    return q;
}

void qdelete(queue_t *q)
{
    pthread_mutex_destroy(&q->lock);
    pthread_cond_destroy(&q->notfull);
    pthread_cond_destroy(&q->notempty);
    free(q);
}
```

# Producer-Consumer (4)

```
void *produce(void *args)
{
    int i, d;
    queue_t *q = (queue_t *)args;
    for (i = 0; i < LOOP; i++) {
        d = rand() % 10;
        put_data(q, d);
        printf("put data %d to queue\n", d);
    }
    pthread_exit(NULL);
}

void *consume(void *args)
{
    int i, d;
    queue_t *q = (queue_t *)args;
    for (i = 0; i < LOOP; i++) {
        d = get_data(q);
        printf("got data %d from queue\n", d);
    }
    pthread_exit(NULL);
}
```

# Exercise

- **void put\_data(queue\_t \*q, int d)**
  - Put data at the end of the queue
  - If the queue is full, wait for the data to be consumed.
- **int get\_data(queue\_t \*q)**
  - Get data in front of the queue
  - If the queue is empty, wait for the data to be produced.